

Problem Set #1

Quiz, 5 questions

✓ **Congratulations! You passed!**

Next Item



1 / 1
point

1.

We are given as input a set of n requests (e.g., for the use of an auditorium), with a known start time s_i and finish time t_i for each request i . Assume that all start and finish times are distinct. Two requests *conflict* if they overlap in time --- if one of them starts between the start and finish times of the other. Our goal is to select a maximum-cardinality subset of the given requests that contains no conflicts. (For example, given three requests consuming the intervals $[0, 3]$, $[2, 5]$, and $[4, 7]$, we want to return the first and third requests.) We aim to design a greedy algorithm for this problem with the following form: At each iteration we select a new request i , including it in the solution-so-far and deleting from future consideration all requests that conflict with i .

Which of the following greedy rules is guaranteed to always compute an optimal solution?

- ☐ At each iteration, pick the remaining request with the earliest start time.
- ☐ At each iteration, pick the remaining request with the fewest number of conflicts with other remaining requests (breaking ties arbitrarily).
- ☒ At each iteration, pick the remaining request with the earliest finish time.

Correct

Let R_j denote the requests with the j earliest finish times. Prove by induction on j that this greedy algorithm selects the maximum-number of non-conflicting requests from S_j .

- ☐ At each iteration, pick the remaining request which requires the least time (i.e., has the smallest value of $t_i - s_i$) (breaking ties arbitrarily).



1 / 1
point

2.

We are given as input a set of n jobs, where job j has a processing time p_j and a deadline d_j . Recall the definition of *completion times* C_j from the video lectures. Given a schedule (i.e., an ordering of the jobs), we define the *lateness* l_j of job j as the amount of time $C_j - d_j$ after its deadline that the job completes, or as 0 if $C_j \leq d_j$. Our goal is to minimize the maximum lateness, $\max_j l_j$.

Which of the following greedy rules produces an ordering that minimizes the maximum lateness? You can assume that all processing times and deadlines are distinct.



Schedule the requests in increasing order of the product $d_j \cdot p_j$

Problem Set #1



None of the other answers are correct.

Quiz, 5 questions



Schedule the requests in increasing order of deadline d_j

Correct

Proof by an exchange argument, analogous to minimizing the weighted sum of completion times.



Schedule the requests in increasing order of processing time p_j



1 / 1
point

3.

In this problem you are given as input a graph $T = (V, E)$ that is a tree (that is, T is undirected, connected, and acyclic). A *perfect matching* of T is a subset $F \subset E$ of edges such that every vertex $v \in V$ is the endpoint of exactly one edge of F . Equivalently, F matches each vertex of T with exactly one other vertex of T . For example, a path graph has a perfect matching if and only if it has an even number of vertices.

Consider the following two algorithms that attempt to decide whether or not a given tree has a perfect matching. The *degree* of a vertex in a graph is the number of edges incident to it. (The two algorithms differ only in the choice of v in line 5.)

Algorithm A:

```

1 While T has at least one vertex:
2   If T has no edges:
3     halt and output "T has no perfect matching."
4   Else:
5     Let v be a vertex of T with maximum degree.
6     Choose an arbitrary edge e incident to v.
7     Delete e and its two endpoints from T.
8 [end of while loop]
9 Halt and output "T has a perfect matching."
```

Algorithm B:

```

1 While T has at least one vertex:
2   If T has no edges:
3     halt and output "T has no perfect matching."
4   Else:
5     Let v be a vertex of T with minimum non-zero degree.
6     Choose an arbitrary edge e incident to v.
7     Delete e and its two endpoints from T.
8 [end of while loop]
9 Halt and output "T has a perfect matching."
```

Is either algorithm correct?



Algorithm A always correctly determines whether or not a given tree graph has a perfect matching; algorithm B does not.



Neither algorithm always correctly determines whether or not a given tree graph has a perfect matching.

Problem Set #1

Quiz, 5 questions

Both algorithms always correctly determine whether or not a given tree graph has a perfect matching.



Algorithm B always correctly determines whether or not a given tree graph has a perfect matching; algorithm A does not.

Correct

Algorithm A can fail, for example, on a three-hop path. Correctness of algorithm B can be proved by induction on the number of vertices in T . Note that the tree property is used to argue that there must be a vertex with degree 1; if there is a perfect matching, it must include the edge incident to this vertex.



1 / 1
point

4.

Consider an undirected graph $G = (V, E)$ where every edge $e \in E$ has a given cost c_e . Assume that all edge costs are positive and distinct. Let T be a minimum spanning tree of G and P a shortest path from the vertex s to the vertex t . Now suppose that the cost of every edge e of G is increased by 1 and becomes $c_e + 1$. Call this new graph G' . Which of the following is true about G' ?



T may not be a minimum spanning tree and P may not be a shortest s - t path.



T may not be a minimum spanning tree but P is always a shortest s - t path.



T is always a minimum spanning tree and P is always a shortest s - t path.



T must be a minimum spanning tree but P may not be a shortest s - t path.

Correct

The positive statement has many proofs (e.g., via the Cut Property). For the negative statement,

think about two different paths from s to t that contain a different number of edges.



1 / 1
point

5.

Suppose T is a minimum spanning tree of the connected graph G . Let H be a connected induced subgraph of G . (i.e., H is obtained from G by taking some subset $S \subseteq V$ of vertices, and taking all edges of E that have both endpoints in S . Also, assume H is connected.) Which of the following is true about the edges of T that lie in H ? You can assume that edge costs are distinct, if you wish. [Choose the strongest true statement.]



For every G and H , these edges are contained in some minimum spanning tree of H

Correct

Proof via the Cut Property (cuts in G correspond to cuts in H with only fewer crossing edges).

Problem Set #1

For every G and H , these edges form a minimum spanning tree of H

Quiz, 5 questions

- ☐ For every G and H and spanning tree T_H of H , at least one of these edges is missing from T_H
 - ☐ For every G and H , these edges form a spanning tree (but not necessary minimum-cost) of H
-

